

Alex Chen

alex.chen.ml@example.com | (415) 892-4512 | linkedin.com/in/alexchen-ml | github.com/alexchen-ml

PROFESSIONAL SUMMARY

Accomplished Senior Machine Learning Engineer with 6+ years of experience designing and deploying deep learning and NLP architectures in cloud-native production environments.

TECHNICAL SKILLS

Python, PyTorch, TensorFlow, Scikit-Learn, Transformers, Hugging Face, NLP, LLM, FastAPI, Docker, Kubernetes, AWS, SQL, CI/CD, Git

WORK EXPERIENCE

Senior Machine Learning Engineer | DeepLogic AI (2021 - Present)

- Designed and implemented transformer-based RAG and search pipelines using PyTorch and Hugging Face, boosting retrieval NDCG@10 by 28%.
- Containerized and deployed high-throughput inference endpoints with FastAPI, Docker, and Kubernetes on AWS EKS.
- Built end-to-end continuous model training pipelines with GitHub Actions and MLflow.

Machine Learning Engineer | DataVision Systems (2018 - 2021)

- Developed production NLP classifiers and recommendation engines using Scikit-Learn and PyTorch.
- Reduced inference latency by 45% via ONNX model quantization.

EDUCATION

Master of Science in Computer Science | Stanford University (2018)

CERTIFICATIONS

AWS Certified Machine Learning - Specialty, TensorFlow Developer Certificate

NOTABLE PROJECTS

Real-Time Semantic Document Search: Built a dense vector retrieval system using Sentence-Transformers, FastAPI, and ChromaDB. Autonomous Code Review Bot: Implemented an automated LLM reviewer using LangChain and GitHub Actions.